

# Erlang assignment 2

## Parallel and distributed programming

Isak Samsten

Spring, 2021

### Introduction

Introductory assignments for parallel Erlang. This week has three assignments to test your knowledge of the concurrency primitives of Erlang. The tasks are awarded a maximum of 20 points. The grades are distributed according to Table 1.

Table 1: Point to grade conversion table

Point	Grade
0 – 8	F
9 – 10	E
11 – 12	D
13 – 15	C
16 – 17	B
18 – 20	A

### Submission

The following files should be submitted to iLearn:

**Problem 1:** `double.erl` and `monitor.erl`

**Problem 2:** `bank.erl`

**Problem 3:** `pmap.erl`

**Problem 4:** `gen_worker.erl`

## Problems

### Problem 1 (3 points)

#### Problem 1.1 (1 points)

Write an Erlang module `double` containing the function `start/0`. The function `start/0` creates a process and registers it under the name `double`. The `double`-process accepts messages on the form `Pid, Ref, N` and returns a message with `N` doubled. If the process receives a non-number it should crash.

#### Example 1.1

```
1> c(double).
2> double:start().
3> Ref = make_ref(), double ! {self(), Ref, 10}.
4> flush().
Shell got {Ref, 20}
5> double ! {self(), Ref, a}.
=ERROR REPORT==== 18-Oct-2019::08:00:13.392934 ===
Error in process <0.85.0> with exit value:
{badarith, [{erlang, '*', [a, 2], []]}}
```

#### Problem 1.2 (2 points)

Write an Erlang module `monitor` containing the function `start/0`. The `start/0` function should start the `double` process and monitor it. If the `double` process crashes, the `monitor` process should restart it and continue to monitor its execution. Finally, implement `double:double/1`, which takes as input an Erlang term and tries to double it using the `double` process, retrying if the `double` process is (temporarily) down due to the `monitor` process restarting it.

#### Example 1.2

```
1> c(double), c(monitor).
2> monitor:start().
3> Ref = make_ref(), double ! {self(), Ref, 10}.
4> flush().
Shell got {..., 20}
5> double ! {self(), Ref, 8}.
6> flush().
Shell got {..., 16}
7> double ! {self(), Ref, a}.
%% the monitor captures the error and restarts the double process
8> 8 = double:double(4).
```

## Problem 2 (7 points)

Write an Erlang module named `bank` (in a file `bank.erl`). The module should have the following functions:

Function	Specification
<code>start()</code>	Start a new bank-server with all accounts having zero balance, returns a <code>pid</code> to the server.
<code>balance(Pid, Who)</code>	Return the balance of <code>Who</code> from the server <code>Pid</code> . Return <code>ok</code> or <code>no_account</code> .
<code>deposit(Pid, Who, X)</code>	Deposit <code>X</code> amount of money to the account <code>Who</code> at the server <code>Pid</code> . If no account exists a new account for <code>Who</code> is opened. Returns <code>{ok, NewAmount}</code> .
<code>withdraw(Pid, Who, X)</code>	Withdraw <code>X</code> amount of money from the account <code>Who</code> at the server <code>Pid</code> . Returns <code>{ok, AmountLeft}</code> or <code>insufficient_funds</code> .
<code>lend(Pid, From, To, X)</code>	Lend <code>X</code> amounts of money from <code>From</code> to <code>To</code> . Return <code>ok</code> , <code>insufficient_funds</code> or <code>{no_account, Who}</code> , where <code>Who</code> is the account that does not exist or the <i>atom</i> both if neither account exists.

Your solution should be implemented using Erlang processes and the implementation should be robust using `make_ref()` and `monitor/2`. If `Pid` has stopped (crashed, or does not exist) the functions should return `no_bank` (e.g., `no_bank = bank:balance(pid(1,2,3), bob)`).

### Example 2

```
1> S = bank:start().
2> no_account = bank:balance(S, bob).
3> {ok, 110} = bank:deposit(S, bob, 110).
4> {ok, 60} = bank:withdraw(S, bob, 50).
5> 60 = bank:balance(S, bob).
6> {no_account, alice} = bank:lend(S, bob, alice, 30).
7> {ok, 0} = bank:deposit(S, alice, 0).
8> ok = bank:lend(S, bob, alice, 30).
9> {ok, 30} = bank:balance(S, bob).
10> {ok, 30} = bank:balance(S, alice).
```

### Problem 3 (5 points)

Implement an Erlang module named `pmap` that provide parallel implementations of the higher-order function `map` (i.e., parallel implementations of the `lists:map(Fun, L)` function). The brave can do Problem 4 first and implement 3.1, 3.2 and 3.3 in terms of it.

#### Problem 3.1 (1 point)

Implement the function `pmap:unordered(Fun, L)` which takes as input a function `Fun` and a list `L`. The function then for each element  $i$  in `L` spawns one *worker* process and evaluates `Fun` on the element. The worker respond to the caller which gather the results and return it. To clarify, the result of `pmap:unordered(Fun, L)` is the same as that of `lists:map(Fun, L)` but without guaranteed order.

##### Example 3.1

```
1> pmap:unordered(fun (X) -> X * 2 end, [1,2,3]).  
[2,4,6] % in undefined order
```

#### Problem 3.2 (2 points)

Implement the function `pmap:unordered(Fun, L, MaxWorkers)` which is functionally equivalent to `pmap:unordered/2` but only spawns `MaxWorkers` *worker*-processes. If `MaxWorkers` is larger than `length(L)`, `MaxWorkers` is set to `length(L)`.

##### Example 3.1

```
1> pmap:unordered(fun (X) -> X * 2 end, [1,2,3], 2).  
[2,4,6] % in undefined order
```

#### Problem 3.3 (2 points)

Implement the function `pmap:ordered(Fun, L, MaxWorkers)` which is functionally equivalent to `pmap:unordered/3` but the order of the returned list is the same as `L`.

##### Example 3.1

```
1> [2, 4, 6] = pmap:ordered(fun (X) -> X * 2 end, [1,2,3], 2).
```

#### Problem 4 (5 points)

Write a behavior-module `gen_worker` that can be used to create work-pools with custom behavior. The `gen_worker` module defines a `handle_work/1-callback`<sup>1</sup> which receives a value and returns `{result, Result}` or `no_result`. The `gen_worker`-module should also handle errors in the user-defined callback function to kill the worker processes. The module also exports the following functions:

Function	Description
<code>Pid = start(Callback, Max)</code>	Start a work-pool with <code>Callback</code> and <code>Max</code> processes handling the work. <code>Pid</code> is the process identifier of the work-pool.
<code>stop(Pid)</code>	Stop <code>Pid</code> and all its workers.
<code>Ref = async(Pid, W)</code>	Schedule <code>W</code> for processing one worker at <code>Pid</code> . <code>Ref</code> is a unique reference that can be used to <code>await(Ref)</code> the result.
<code>await(Ref)</code>	Await the result with the unique reference <code>Ref</code> created by <code>async(Pid, W)</code> . Returns <code>no_result</code> , <code>error</code> or <code>{result, Result}</code> .
<code>await_all(Refs)</code>	Await the work for all references in the list <code>Refs</code> and return a (possibly empty) list of results. Work resulting in <code>no_result</code> should <i>not</i> be included in the list.

More specifically, `gen_worker` encapsulate the behavior of a work-pool, i.e., it defines a *master* process that receives and distributes work to *k worker* processes. The worker processes in turn call `Callback:handle_work(Work)` and sends back the result to the *master* process. The *master* process then sends back the result to the caller of `gen_worker:async/2` (note that `async/2` does not wait for the result but return a reference that can be used by `await/1` to wait for the result). The caller can then use `gen_worker:await/1` to receive the result.

---

<sup>1</sup>Implementors should receive a warning if the callback is not exported, i.e., you should use `-callback` to define your callback function.

On the next page, we have a small module `pmap` which uses the `gen_worker` module to implement `pmap:ordered/2`, which can be used as in Example 4.

#### Example 4

```
1> c(gen_worker), c(pmap).
2> pmap:ordered(fun (X) -> X*2 end, [1,2,3]).
[2, 4, 6]
```

#### `pmap.erl`

```
-module(pmap).
-behaviour(gen_worker).
-export([handle_work/1, ordered/2]).

%% Simple handle_work: apply the function to the value
handle_work({Fun, V}) ->
    {result, Fun(V)}.

ordered(Fun, L) ->
    %% Start a work-pool with 2 workers
    WorkPool = gen_worker:start(?MODULE, 2),

    %% Schedule the work asynchronously
    Refs = [gen_worker:async(WorkPool, {Fun, V}) || V <- L],

    %% Await the result
    Result = gen_worker:await_all(Refs),

    %% Stop our work pool
    gen_worker:stop(S),

    %% Return the result
    Result.
```